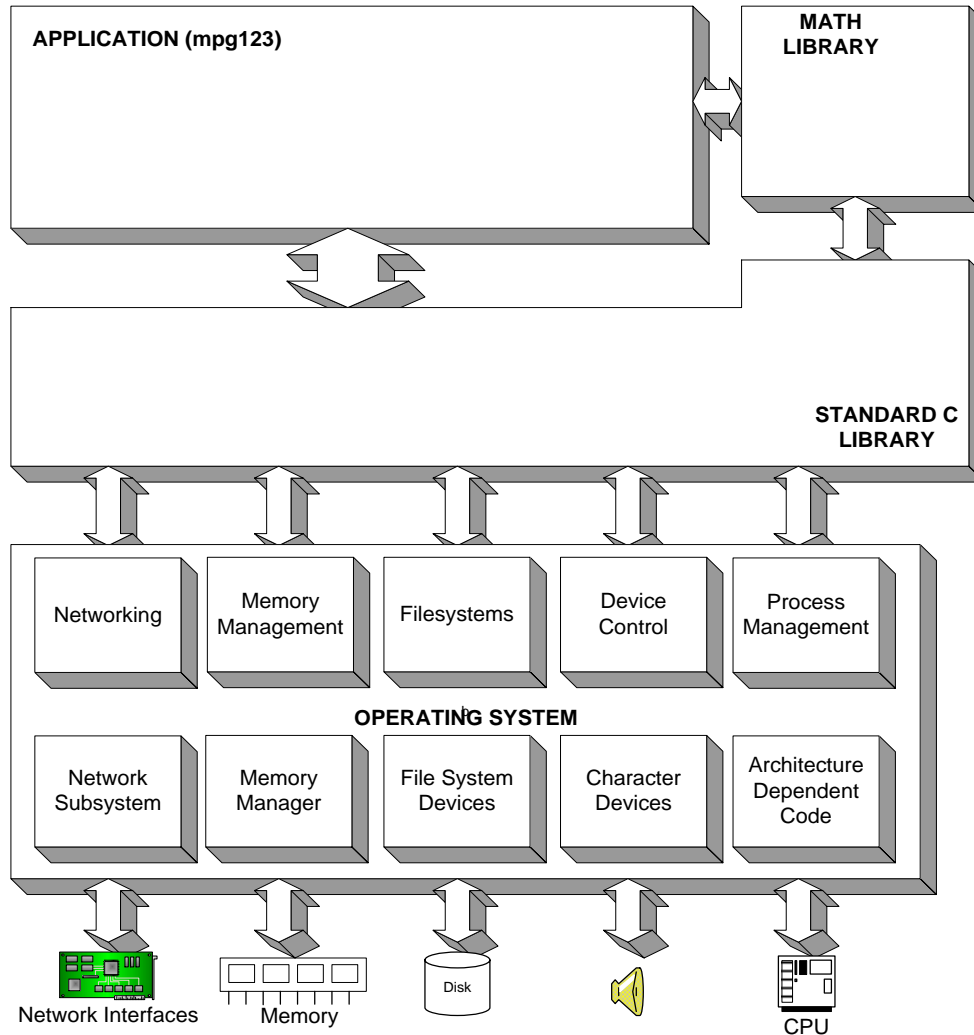


Kernels and Device Drivers

Linux Execution Environment



- Program
- Libraries
- Kernel subsystems

Process

- Process: program in execution.
 - Unique “pid”. Hierarchy
- User address space vs. kernel address space
- Application requests OS services through TRAP mechanism
 - x86: syscall number in eax register, exception (int \$0x80)
 - result = read (file descriptor, user buffer, amount in bytes)
 - Read returns real amount of bytes transferred or error code (<0)

- Kernel has access to kernel address space (code, data, and device ports and memory), and to user address space, but only to the process that is currently running
- “**Current**” process descriptor. “**current**→**pid**” points to current pid
- Two stacks per process: user stack and kernel stack
- Special instructions to copy parameters / results between user and kernel space

Splitting the Kernel

- Process management
 - Creates, destroys processes
 - Supports communication among processes
 - Signals, pipes, etc.
 - Schedules how processes share the CPU
- Memory management
 - Virtual addressing

Splitting the Kernel

- File systems
 - Everything in UNIX can be treated as a file
 - Linux supports multiple file systems
- Device control
 - Every system operation maps to a physical device
 - Few exceptions: CPU, memory, etc.
- Networking
 - Handles packets
 - Handles routing and network address resolution issues

Device drivers

- Device drivers
 - Black boxes to hide details of hardware devices
 - Use standardized calls
 - Independent of the specific driver
 - Main role
 - Map standard calls to device-specific operations
 - Can be developed separately from the rest of the kernel
 - Plugged in at runtime when needed

The Role of the Device Driver

- Implements the *mechanisms* to access the hardware
 - E.g., show a disk as an array of data blocks
- Does not force particular *policies* on the user
 - Examples
 - Who many access the drive
 - Whether the drive is accessed via a file system
 - Whether users may mount file systems on the drive

Policy-Free Drivers

- A common practice
 - Support for synchronous/asynchronous operation
 - Be opened multiple times
 - Exploit the full capabilities of the hardware
- Easier user model
- Easier to write and maintain
- To assist users with policies, release device drivers with user programs

Loadable Modules

- The ability to add and remove kernel features at runtime
- Each unit of extension is called a *module*
- Use **insmod** program to add a kernel module
- Use **rmmmod** program to remove a kernel module

Classes of Devices and Modules

- Character devices
- Block devices
- Network devices
- Others

Character Devices

- Abstraction: a stream of bytes
 - Examples
 - Text console (`/dev/console`)
 - Serial ports (`/dev/ttyS0`)
 - Usually supports **open**, **close**, **read**, **write**
 - Accessed sequentially (in most cases)
 - Might not support file seeks
 - Exception: frame grabbers
 - Can access acquired image using **mmap** or **lseek**

Block Devices

- Abstraction: array of storage blocks
- However, applications can access a block device in bytes
 - Block and char devices differ only at the kernel level
 - A block device can host a file system

Network Devices

- Abstraction: data packets
- Send and receive packets
 - Do not know about individual connections
- Have unique names (e.g., **eth0**)
 - Not in the file system
 - Support protocols and streams related to packet transmission (i.e., no **read** and **write**)

Other Classes of Devices

- Examples that do not fit to previous categories:
 - USB
 - SCSI
 - FireWire
 - MTD

File System Modules

- Software drivers, not device drivers
- Serve as a layer between user API and block devices
- Intended to be device-independent