

System Architecture and Assembly

Systems Programming (CST-210)

A. P. MAZUMDAR

2

Intel x86 Processors

- ▶ Totally dominate laptop/desktop/server market
- ▶ Evolutionary design
 - ▶ Backwards compatible up until 8086, introduced in 1978
 - ▶ Added more features as time goes on

Intel x86 Processors

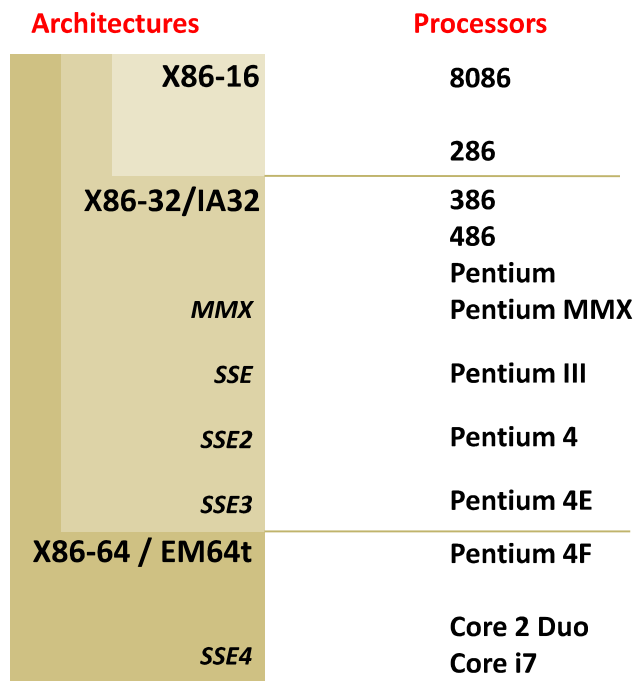
- ▶ Complex instruction set computer (CISC)
 - ▶ Many different instructions with many different formats
 - ▶ But, only small subset encountered with Linux programs
 - ▶ Hard to match performance of Reduced Instruction Set Computers (RISC)
 - ▶ But, Intel has done just that!
 - ▶ In terms of speed. Less so for low power.

Intel x86 Evolution: Milestones

| <i>Name</i> | <i>Date</i> | <i>Transistors</i> | <i>MHz</i> |
|--|-------------|--------------------|--------------|
| ▶ 8086 | 1978 | 29K | 5-10 |
| <ul style="list-style-type: none"> ▶ First 16-bit processor. Basis for IBM PC & DOS ▶ 1MB address space | | | |
| ▶ 386 | 1985 | 275K | 16-33 |
| <ul style="list-style-type: none"> ▶ First 32 bit processor , referred to as IA32 ▶ Added “flat addressing” ▶ Capable of running Unix ▶ 32-bit Linux/gcc uses no instructions introduced in later models | | | |

Intel x86 Evolution: Milestones

| <i>Name</i> | <i>Date</i> | <i>Transistors</i> | <i>MHz</i> |
|---|-------------|--------------------|------------------|
| ► Pentium 4F | 2004 | 125M | 2800-3800 |
| ► First 64-bit processor, referred to as x86-64 | | | |
| ► Core i7 | 2008 | 731M | 2667-3333 |
| ► New machines | | | |



IA: often redefined as latest Intel architecture

x86 Clones: Advanced Micro Devices (AMD)

► Historically

- AMD has followed just behind Intel
- A little bit slower, a lot cheaper

► Then

- Recruited top circuit designers from Digital Equipment Corp. and other downward trending companies
- Built Opteron: tough competitor to Pentium 4
- Developed x86-64, their own extension to 64 bits

Intel's 64-Bit

- Intel Attempted Radical Shift from IA32 to IA64
 - Totally different architecture (Itanium)
 - Executes IA32 code only as legacy
 - Performance disappointing
- AMD Stepped in with Evolutionary Solution
 - x86-64 (now called "AMD64")
- Intel Felt Obligated to Focus on IA64
 - Hard to admit mistake or that AMD is better

Intel's 64-Bit

- ▶ 2004: Intel Announces EM64T extension to IA32
 - ▶ Extended Memory 64-bit Technology
 - ▶ Almost identical to x86-64!
- ▶ All but low-end x86 processors support x86-64
 - ▶ But, lots of code still runs in 32-bit mode

IA32 (Pentium) Processor Architecture

- ▶ **32 bit Processor**
- ▶ **1 WORD = 16bit**
- ▶ **32 bits = 2 WORDS = 1 Double Word (DWORD)**

Processor modes

1. Protected (*important*)
 - ▶ 32-bit mode
 - ▶ 32-bit (4GB) address space
2. Virtual 8086 modes
3. Real mode
 - ▶ 1MB address space
4. System management mode

Registers

- ▶ 32-bit GPR's ("general" purpose registers):

| | |
|---------------|------------|
| eax | ebp |
| ebx | esp |
| ecx | esi |
| edx | edi |
| eflags | eip |

e[a,b,c,d]x:



Note: eax is one register that can be viewed four different ways.

Registers

- ▶ Not really GPR's.
 - ▶ eax - accumulator; multiplication and division
 - ▶ ecx - loop counter
 - ▶ esp - stack pointers; don't use
 - ▶ esi, edi - for memory-to-memory transfer
 - ▶ ebp - used by HLL for local vars on stack

Registers

- ▶ Additional registers:
 - ▶ 16-bit segment registers
 - ▶ cs, es, ss, fs, ds, gs
 - ▶ don't use
 - ▶ eip
 - ▶ instruction pointer / program counter (PC)
 - ▶ don't use

Registers

- ▶ CS
 - ▶ Address of current code segment
- ▶ SS
 - ▶ Address of current stack segment
- ▶ Others (DS, ES, FS, GS)
 - ▶ Address of data segments

Registers

- ▶ Additional registers:
 - ▶ eflags
 - ▶ contains results of operations
 - ▶ 32 individual bits
 - ▶ control flags
 - ▶ status flags:
 - ▶ **C** = carry (unsigned)
 - ▶ **O** = overflow (signed); also called **V**
 - ▶ **S** = sign; also called **N** for negative
 - ▶ **Z** = zero

Registers

- ▶ Additional registers:
 - ▶ floating point registers:
 - ▶ ST(0) ... ST(7)
 - ▶ 80 bits
 - ▶ MMX has 8 64-bit regs
 - ▶ Translate segment address to Physical address
 - ▶ XMM has 8 128-bit regs

Compilation

- ▶ Two parts of a program: **p1.c** and **p2.c**
- ▶ To compile:
 - ▶ `gcc -O1 -o p p1.c p2.c`
- ▶ `-o`
 - ▶ Output file name, followed by the <name>
- ▶ `-O1`
 - ▶ Level of optimization: (higher level = faster execution, slower compilation)

Compilation

- ▶ ASM code generation:
 - ▶ **`gcc -O1 -S code.c`**
- ▶ Output file `code.s` will be generated
- ▶ `-S` : gcc option to compile till assembly level

```

1  int accum = 0;
2
3  int sum(int x, int y)
4  {
5      int t = x + y;
6      accum += t;
7      return t;
8  }

sum:
    pushl   %ebp
    movl    %esp, %ebp
    movl    12(%ebp), %eax
    addl    8(%ebp), %eax
    addl    %eax, accum
    popl    %ebp
    ret

```

Creating Object code

► gcc -O1 -c code.c

- **-C:** option for generating obj code

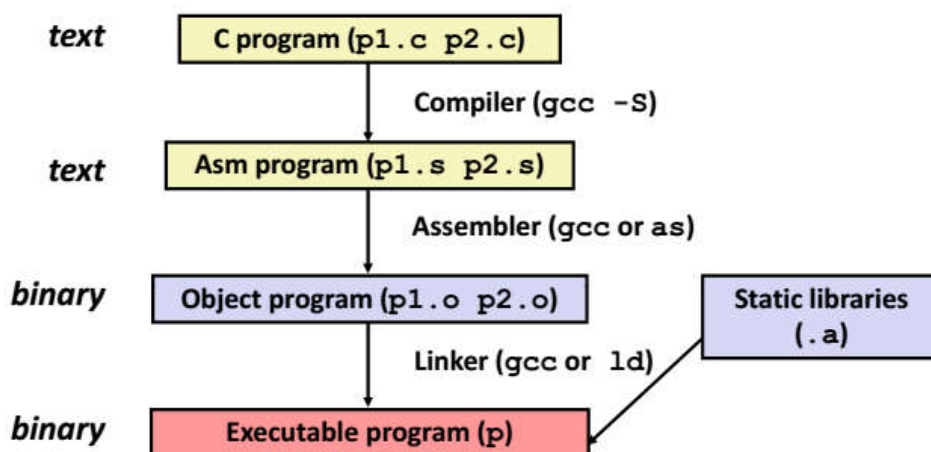
► For code.c :

```
55 89 e5 8b 45 0c 03 45 08 01 05 00 00 00 00 5d c3
```

Generating Object code

- ▶ **gcc -O1 -c code.c**
- ▶ Use Disassembler to byte code length
 - ▶ In this case it is 17
- ▶ Use GNU debugging tool (GDB) on **code.o** to get the code
 - ▶ (gdb) x/17xb sum

Summary of compilation



Revisit IA32 Integer Registers

| 31 | 15 | 8 | 7 | 0 |
|------|-----|-----|-----|---|
| %eax | %ax | %ah | %al | |
| %ecx | %cx | %ch | %cl | |
| %edx | %dx | %dh | %dl | |
| %ebx | %bx | %bh | %bl | |
| %esi | %si | | | |
| %edi | %di | | | |
| %esp | %sp | | | |
| %ebp | %bp | | | |

Stack pointer

Frame pointer