

MONITORS

Dr. Emmanuel S. Pilli

Monitors

- Semaphore is a low level primitive because its is unstructured
- In case of large systems being built, using semaphores alone, the responsibility for the correct use of semaphores would be diffused among all the implementations of the system
- Monitors provide a structured concurrent programming primitive that concentrates the responsibility for correctness into modules

Monitors

- Monitors are generalizations of *kernel* or *supervisor* found in operating systems, where critical sections such as the allocation of memory are centralized in a privileged program
- Monitors are decentralized versions of the monolithic kernels
- A separate monitor is defined for each object or related group of objects that requires synchronization

Monitors

- If operations of the same monitor are called by more than one process, the implementation ensures that these are executed under mutual exclusion
- If operations of different monitors are called, their executions can be interleaved
- Monitors have become an extremely important synchronization mechanism because they are a natural generalization of the *object* of OOP, which encapsulates data and operation declarations within a *class*

Monitors

- Objects of a class can be allocated at runtime and the operations of the class invoked on the fields of the object
- Monitor adds the requirement that only one process can execute an operation on an object at any one time

Monitors

- While the fields of the object may be declared either *public* (directly accessible outside the class) or *private* (accessible only by operations declared within the class), the fields of a monitor are all *private*.
- This ensures that the fields of a monitor are accessed consistently
- Implementations of monitors in programming languages are quite different from one another

Monitors

Algorithm 7.1: Atomicity of monitor operations

monitor CS

integer $n \leftarrow 0$

operation increment

integer temp

temp $\leftarrow n$

$n \leftarrow \text{temp} + 1$

p

p1: CS.increment

q

q1: CS.increment


Monitors

- Monitor CS contains one variable **n** and one operation **increment**
- Two statements are contained within this operation, together with the declaration of a local variable
- The variable **n** is not accessible outside the monitor
- Two processes **p** and **q**, each call the monitor operation **CS.increment**

Monitors

- Only one process can execute the monitor operation and as mutual exclusion is ensured in access to the variable
- As one process is executing the statements of a monitor operation, the other statements are waiting outside
- This also solves the critical section problem as a critical section statements can be placed within the operation of a monitor

Monitors



Monitors

- The statements of critical section are encapsulated in the monitor rather than replicated in each process
- The synchronization is implicit and does not require the wait and signal statements
- Monitor is a static entry and not a dynamic process ... It is a set of operations that just sit there (keep waiting) waiting for a process to invoke one of them

Monitors

- There is an implicit lock on the door to the monitor, ensuring only one process is inside the monitor at any time
- As with semaphores, if there are several processes attempting to enter a monitor, only one of them will succeed
- There is no explicit queue associated with the monitor entry, so starvation is possible

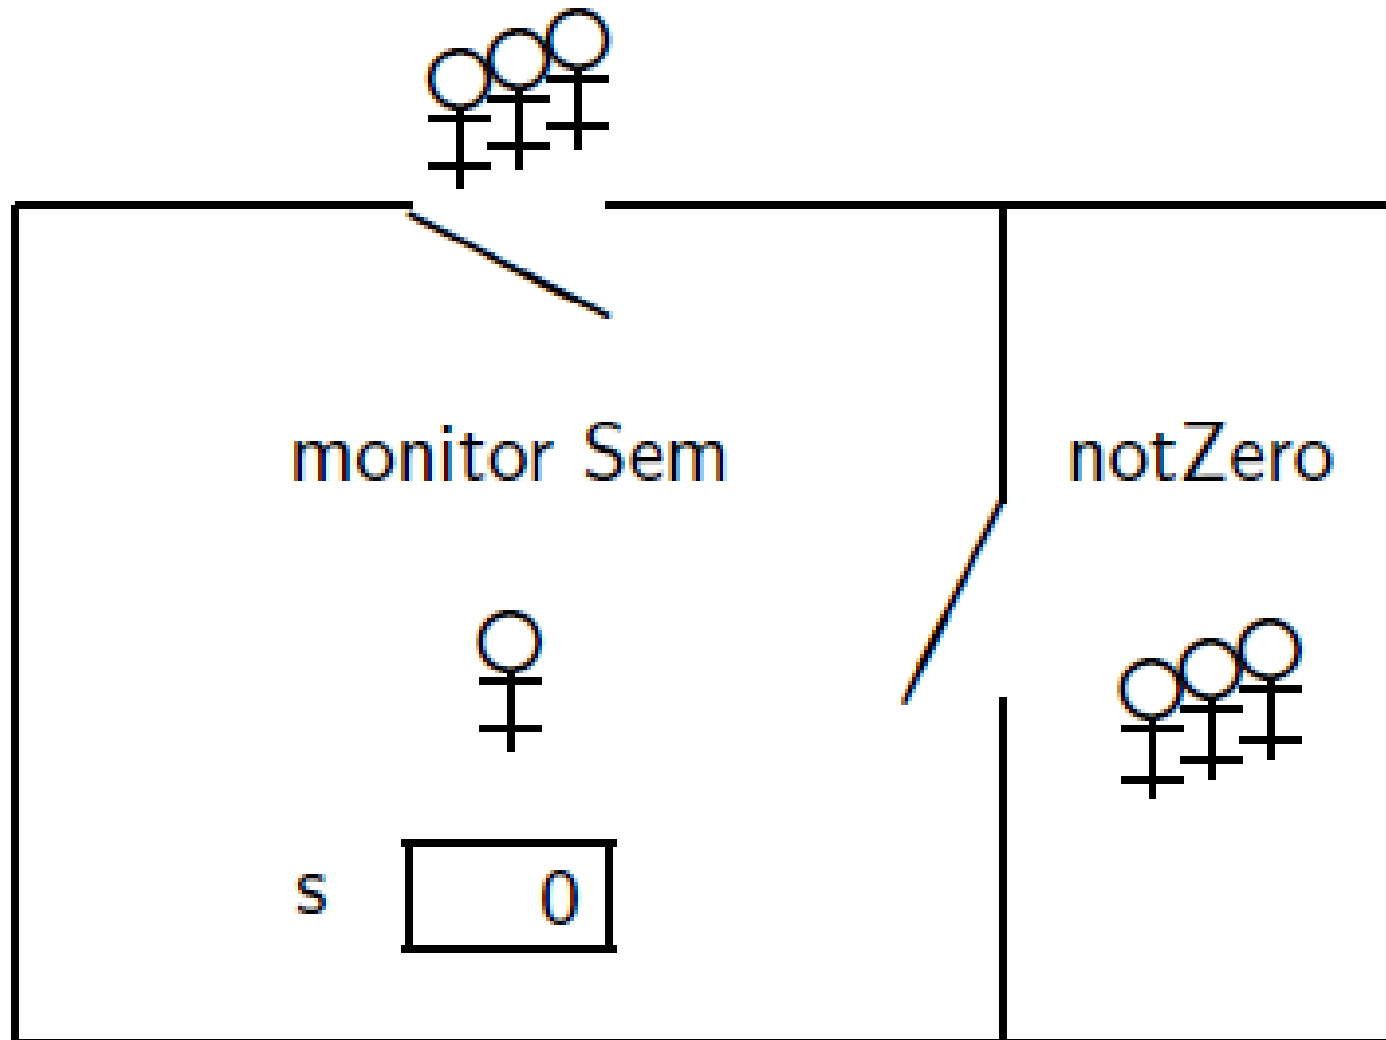
Condition Variables

- Monitor implicitly enforces mutual exclusion to its variables
- Problems in concurrent programming have explicit synchronization requirements
- Two approaches in providing synchronization:
 - Required condition is named by an explicit *condition variable* (called an *event*). Ordinary boolean expressions are used to test the condition
 - Block directly on the expression and let the implementation implicitly unblock a process when the expression is true

Condition Variables

- Condition variables provide a mechanism to wait for events (a “rendezvous point”)
 - Resource available, no more writers, etc.
- Condition variables support three operations:
 - Wait – release monitor lock, wait for C/V to be signaled
 - So condition variables have wait queues, too
 - Signal – wakeup one waiting thread
 - Broadcast – wakeup all waiting threads

Simulating Semaphores



Algorithm 7.2: Semaphore simulated with a monitor

monitor Sem

integer $s \leftarrow k$

condition notZero

operation wait

if $s = 0$

waitC(notZero)

$s \leftarrow s - 1$

operation signal

$s \leftarrow s + 1$

signalC(notZero)

p

loop forever

non-critical section

p1: Sem.wait

critical section

p2: Sem.signal

q

loop forever

non-critical section

q1: Sem.wait

critical section

q2: Sem.signal

Simulating Semaphores

- Integer component of semaphore is stored in variable `s`
- Condition variable `cond` implements queue of blocked processes
- Condition variables are named with the condition you want to be true
- `waitC(notZero)` is read as “wait for notZero to be true” and `signalC(notZero)` is read as “signal that notZero is true”
- If the value of `s` is zero, the process executing `Sem.wait` executes the monitor statement `waitC(notZero)`. The process is said to be blocked on the condition

Simulating Semaphores

- Process executing waitC blocks unconditionally as we assume that the condition has been tested for in the preceding if statement
- As monitor operation is atomic the value of the condition cannot change between testing its value and executing waitC
- When a process executes a Sem.signal operation, it unblocks the first process blocked on that condition
- With each condition variable is associated a FIFO queue of blocked processes

Operations on Condition Variables

waitC(cond)

append p to cond
p.state \leftarrow blocked
monitor.lock \leftarrow release

empty(cond)

return cond = empty

signalC(cond)

if cond \neq empty
 remove head of
 cond and assign to q
 q.state \leftarrow ready

Semaphore vs Monitor

Semaphore

- Wait may or may not block
- Signal always has an effect
- Signal unblocks an arbitrary blocked process
- A process unblocked by signal can resume execution immediately

Monitor

- WaitC always blocks
- SignalC has no effect if queue is empty
- SignalC unblocks the process at the head of the queue
- A process unblocked by signalC must wait for the signalling process to leave monitor

Algorithm 7.3: Producer-consumer (finite buffer, monitor)

monitor PC

bufferType buffer \leftarrow empty

condition notEmpty

condition notFull

operation append(datatype V)

if buffer is full

waitC(notFull)

append(V, buffer)

signalC(notEmpty)

operation take()

datatype W

if buffer is empty

waitC(notEmpty)

W \leftarrow head(buffer)

signalC(notFull)

return W

Producer Consumer

Algorithm 7.3: Producer-consumer (finite buffer, monitor) (continued)

producer	consumer
<pre>datatype D loop forever p1: D ← produce p2: PC.append(D)</pre>	<pre>datatype D loop forever q1: D ← PC.take q2: consume(D)</pre>

Emulating Semaphores using Monitors

monitor **Semaphore_Emulation** is

S: Integer := S0;
Not_Zero : Condition;

Procedure **Semaphore_Wait** is
begin
 if S=0 then Wait(Not_Zero); end if;
 S := S-1;
end **Semaphore_Wait**;

Procedure **Semaphore_Signal** is
begin
 S := S+1; Signal(Not_Zero);
end **Semaphore_Signal**;

end monitor

Readers Writers Problem

- Similar to Mutex problem where several processes are competing for access to a critical section
- Readers are processes which require to exclude writers but not other readers
- Writers are processes which require to exclude both readers and other writers
- Several processes can read data concurrently, but writing or modifying data must be done under Mutex to ensure consistency of Data

Algorithm 7.4: Readers and writers with a monitor

monitor RW

integer readers $\leftarrow 0$

integer writers $\leftarrow 0$

condition OKtoRead, OKtoWrite

operation StartRead

if writers $\neq 0$ or not empty(OKtoWrite)

waitC(OKtoRead)

readers \leftarrow readers + 1

signalC(OKtoRead)

operation EndRead

readers \leftarrow readers - 1

if readers = 0

signalC(OKtoWrite)

Algorithm 7.4: Readers and writers with a monitor (continued)

operation StartWrite

if writers \neq 0 or readers \neq 0

waitC(OKtoWrite)

writers \leftarrow writers + 1

operation EndWrite

writers \leftarrow writers - 1

if empty(OKtoRead)

then signalC(OKtoWrite)

else signalC(OKtoRead)

reader	writer
p1: RW.StartRead	q1: RW.StartWrite
p2: read the database	q2: write to the database
p3: RW.EndRead	q3: RW.EndWrite

Readers Writers Problem

- Monitor uses 4 variables
 - Readers: The number of threads currently reading the database after successfully executing StartRead but before executing EndRead
 - writers: The number of writers currently writing to the database after successfully executing StartWrite but before executing EndWrite
 - OktoRead: A condition variable for blocking readers until it is “OK to read”
 - OktoWrite: A condition variable for blocking writers until it is “OK to write”

Dining Philosophers Problem

Algorithm 7.5: Dining philosophers with a monitor

```
monitor ForkMonitor
  integer array[0..4] fork  $\leftarrow$  [2, ..., 2]
  condition array[0..4] OKtoEat

  operation takeForks(integer i)
    if fork[i]  $\neq$  2
      waitC(OKtoEat[i])
    fork[i+1]  $\leftarrow$  fork[i+1] - 1
    fork[i-1]  $\leftarrow$  fork[i-1] - 1

  operation releaseForks(integer i)
    fork[i+1]  $\leftarrow$  fork[i+1] + 1
    fork[i-1]  $\leftarrow$  fork[i-1] + 1
    if fork[i+1] = 2
      signalC(OKtoEat[i+1])
    if fork[i-1] = 2
      signalC(OKtoEat[i-1])
```

Dining Philosophers Problem

Algorithm 7.5: Dining philosophers with a monitor (continued)
philosopher i
loop forever
p1: think
p2: takeForks(i)
p3: eat
p4: releaseForks(i)

Dining Philosophers Problem

- Monitor maintains an array of fork which counts the number of free forks available to each philosopher
- The takeForks operation waits on a condition variable until two forks are available
- It decrements the number of forks available to its neighbor before leaving the monitor
- After eating, a philosopher calls releaseForks which updates the array fork and checks if freeing these forks makes it possible to signal

What's wrong with Semaphores

- They are essentially shared global variables.
- There is no linguistic connection between the semaphore and the data to which the semaphore controls access.
- Access to semaphores can come from anywhere in a program.
- They serve two purposes, mutual exclusion and scheduling constraints.
- There is no control or guarantee of proper usage.

Monitor – A Formal Definition

- A Monitor defines a lock and zero or more condition variables for managing concurrent access to shared data.
- The monitor uses the lock to insure that only a single thread is active in the monitor at any instance.
- The lock also provides mutual exclusion for shared data.
- Condition variables enable threads to go to sleep inside of critical sections, by releasing their lock at the same time it puts the thread to sleep.

Monitor Operations

- Encapsulates the shared data you want to protect.
- Acquires the mutex at the start.
- Operates on the shared data.
- Temporarily releases the mutex if it can't complete.
- Reacquires the mutex when it can continue.
- Releases the mutex at the end.